

# Statistical Computing

Michael Mayer

February 2023



# Statistical Computing: What will we do?

## Chapters

1. R in Action
2. Statistical Inference
3. Linear Models
4. Model Selection and Validation
5. Trees
6. Neural Nets

## Remarks

- ▶ Chapters 3 to 6:  
Statistical ML in Action
- ▶ Two weeks per chapter
- ▶ Exercises at end of chapter notes

# Neural Nets

# Outline

- ▶ Understanding Neural Nets
- ▶ Practical Considerations
- ▶ Extended examples

# Neural Nets

- ▶ Around since the 1950ies
- ▶ Underwent different development steps, e.g.
  - ▶ use of backpropagation
  - ▶ GPUs
- ▶ Black Box
- ▶ TensorFlow/Keras, PyTorch

## “Swiss Army Knife” among ML Algorithms

**Can fit linear models**

**Learn interactions  
and non-linear terms**

**>1 Responses possible**

**Flexible and mixed  
in- and output dimensions**

**Fit data larger than RAM**

**Non-linear**      **Learn «online»**  
**dimension reduction**

**Sequential and spatial  
in- and output**

**Flexible loss functions**

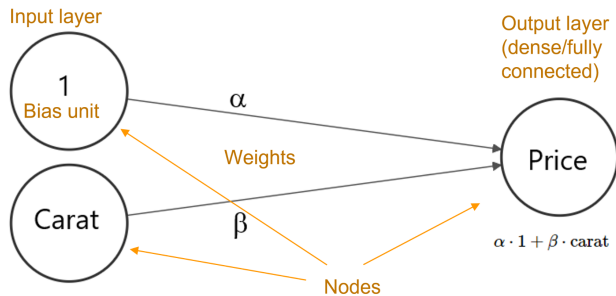
# Understanding Neural Nets in three Steps

1. Linear regression as neural net
2. Hidden layers
3. Activation functions

Using **diamonds** data

## Step 1: Linear Regression as Neural Net

- ▶  $\mathbb{E}(\text{price}) = \alpha + \beta \cdot \text{carat}$
- ▶ OLS  
 $\hat{\alpha} \approx -2256, \hat{\beta} \approx 7756$
- ▶ Represented as neural network graph



Example



# The Optimization Algorithm

## Mini-batch gradient descent with backpropagation

Notation: Neural net  $f_\beta$ ; its total loss on data  $D$  and loss function  $L$ :

$$Q(f_\beta, D) = \sum_{(y_i, \mathbf{x}_i) \in D} L(y_i, f_\beta(\mathbf{x}_i))$$

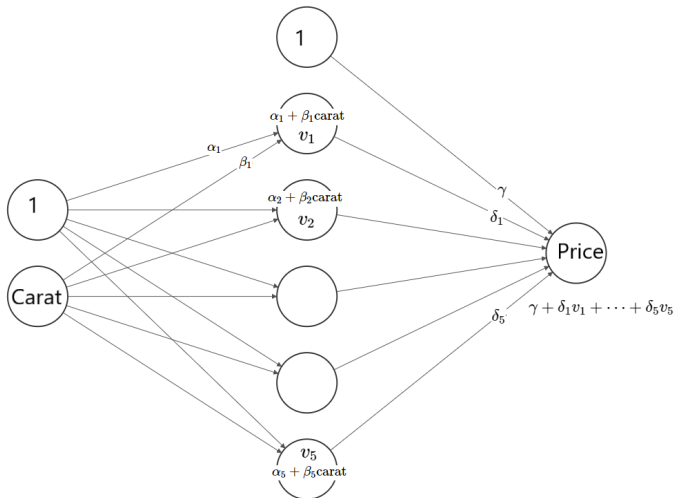
1. Init: Randomly initialize parameter vector  $\beta$  by  $\hat{\beta}$
2. Forward: Calculate  $Q(f_{\hat{\beta}}, D_{\text{batch}})$  on **batch**
3. Backprop: Modify  $\hat{\beta}$  to improve  $Q(f_{\hat{\beta}}, D_{\text{batch}})$ 
  - 3.1 Calculate partial derivatives  $\nabla \hat{\beta} = \frac{\partial Q(f_{\hat{\beta}}, D_{\text{batch}})}{\partial \beta} \big|_{\beta=\hat{\beta}}$  using backprop (=?)
  - 3.2 Gradient descent: Move slightly into right direction:  $\hat{\beta} \leftarrow \hat{\beta} - \lambda \cdot \nabla \hat{\beta}$
4. Repeat Steps 2 and 3 until one **epoch** is over
5. Repeat Step 4 until some stopping criterion triggers

SGD? Local minima?

## Step 2: Hidden Layers

- ▶ Add **hidden layers** for more parameters (= flexibility)
- ▶ Their nodes are latent/implicit variables
- ▶ Representational learning
- ▶ **Encoding?**
- ▶ **Deep** neural net?

### Example



## Step 3: Activation Functions

Non-linear transformations  $\sigma$  of node values necessary!

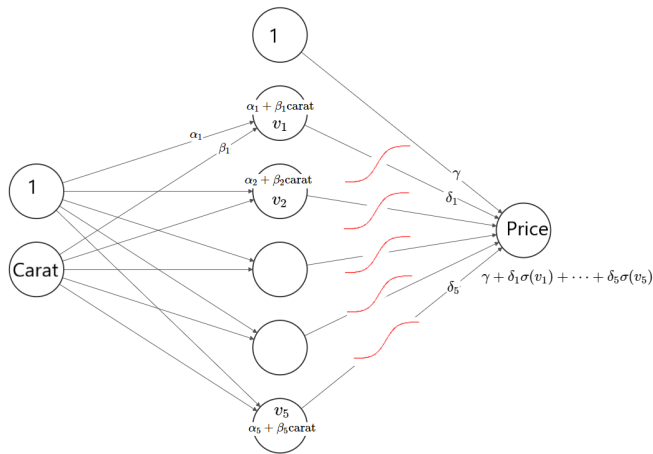
- ▶ tanh:  $\sigma(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$
- ▶ ReLU:  $\sigma(x) = \max(0, x)$



### Two purposes

- ▶ Imply interactions and non-linear terms
- ▶ Inverse link as in GLMs

### Example



# Practical Considerations

**Validation and tuning  
of main parameters**

**Callbacks**

**Overfitting and  
regularization**

**Input standardization**

**Missing values**

**Types of layers**

**Optimizer**

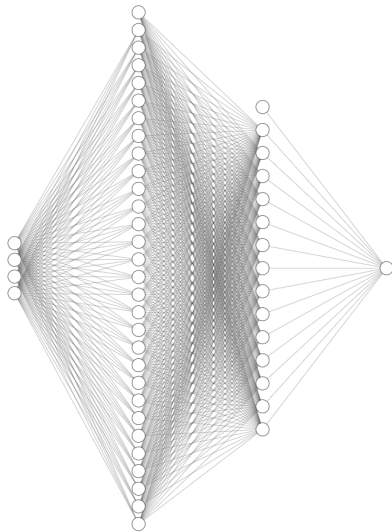
**Choosing the  
architecture**

**Categorical input**

**Interpretation**

**Custom losses and  
evaluation metrics**

## Example: Diamonds



## Excursion: Model-Agnostic Importance Measure

Permutation importance of feature  $X^{(j)}$ , data  $D$ , and performance measure  $S$ :

$$\text{PVI}(j, D) = S(\hat{f}, D^{(j)}) - S(\hat{f}, D)$$

- ▶  $D^{(j)}$  is version of  $D$  with randomly permuted values in  $j$ -th feature column
- ▶ Read: How much  $S$  worsens after shuffling column  $j$ ?  
The larger, the more important. If 0, feature is unimportant
- ▶ Computationally cheap  $\rightarrow$  repeat  $m$  times
- ▶ Model is never refitted
- ▶ Training or test data?

### Example

# Embeddings

Represent unordered categorical  $X$  with  $K$  levels by  $m \ll K$  numeric features

## Embedding layer

- ▶  $X$  integer encoded
- ▶ Dummy matrix  $\tilde{X}$  with  $K$  columns
- ▶ Multiply  $\tilde{X}$  with  $(K \times m)$  matrix  $\beta$
- ▶ Embedding matrix  $\beta$  estimated like other parameters
- ▶ Trick:  $\tilde{X}\beta$  is calculated via index slicing from  $X$  and  $\beta$   
→  $\tilde{X}$  is never materialized

## Example

Taxi trips

## Excursion: Analysis Scheme X

$T(Y)$ : quantity of interest

### Steps

1. Calculate  $T(Y)$  on the full data
2. Calculate  $T(Y)$  stratified by covariates  $X^{(j)} \rightarrow$  bivariate associations
3. Accompany Step 2 by ML model  $\rightarrow$  multivariate associations
  - ▶ Study model performance
  - ▶ Study variable importance  $\rightarrow$  sort results of Step 2
  - ▶ Study PDP (or similar) for each  $X^{(j)}$  and compare with Step 2

### Example



# Comparison of ML Algorithms

Aspect	GLM	Neural Net	Decision Tree	Boosting	Random Forest	k-Nearest Neighbour
Scalable	😍	😍	😊	😊	😐	😞
Easy to tune	😐	😐	😐	😐	😊	😐
Flexible losses	😊	😍	😊	😊	😐	😐
Regularization	✓	✓	✓	✓	✓	✓
Case weights	✓	✓	✓	✓	✓	✓
Missing input allowed	😞	😞	✓	✓	😞	😞
Interpretation	😍	😐	😍	😐	😐	😐
Space on disk	😍	😍	😍	😊	😞	😞
Birth date (approx.)	1972 (Nelder & Wedderburn)	1974 Backprop (Werbos)	1984 (Breiman et al.)	1990 (Schapire)	2001 (Breiman)	1951 (Fix & Hodges)